

# Optimizing Write-Heavy Database Operations Using $B^\epsilon$ -Trees

---


Christoph Rotte

September 1, 2022

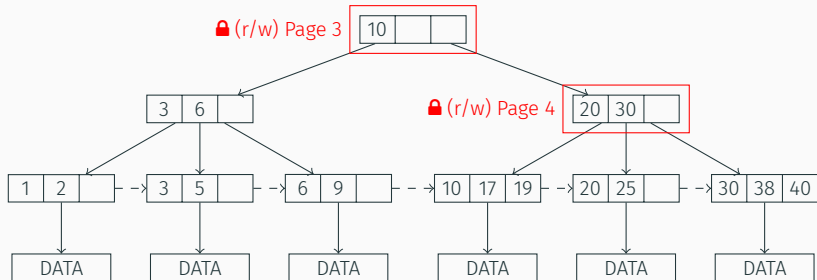
Bachelor's Thesis - Final Presentation

Chair for Database Systems | Technical University of Munich


# B/B<sup>+</sup>-Trees in Databases



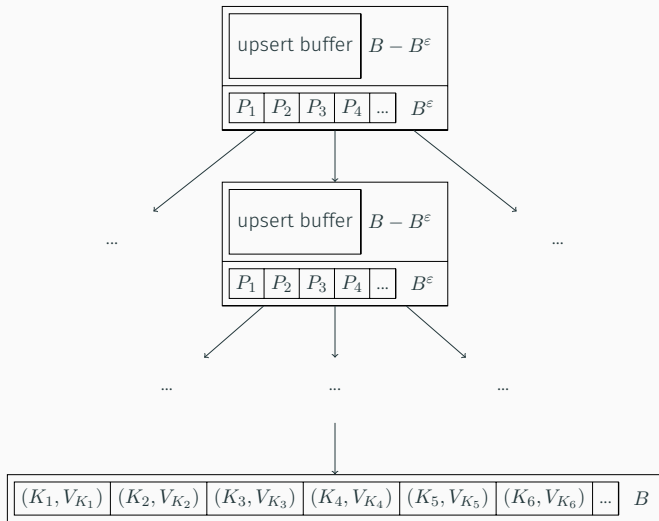
# B/B<sup>+</sup>-Trees in Databases




# Log-Structured Merge-Trees (LSM-Trees)



# $B^\epsilon$ -Trees | Structure




# $B^\varepsilon$ -Trees | Lookups



## Asymptotic Comparison | I/O Operations


	B/B <sup>+</sup> -Tree	B <sup>ε</sup> -Tree	LSM-Tree
Upserts	-	+	+
Lookups	+	+	-

# Implementation | Preemptive Splitting

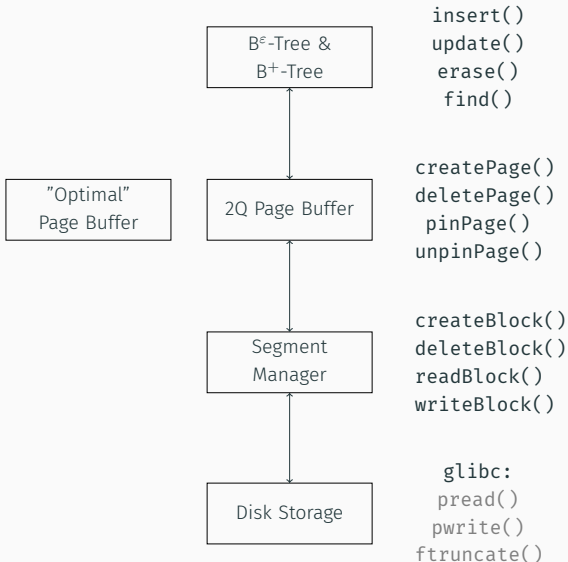




# Implementation | Separate Root Node



# Implementation | Design Layers



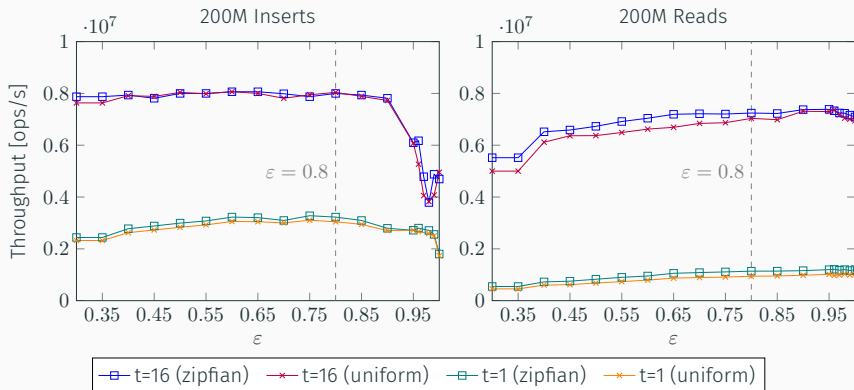
- **Benchmark Tool:** Unum Cloud Benchmark (UCSB)  
→ rewrite of YCSB in C++
- Intel i9-7900X (20 threads) | Samsung SSD 970 EVO
- **Sizes:** 16KiB pages | 8B keys | 100B values
- **Distribution:** Zipfian (constant: 0.99)

# "Optimal" Buffer | 200M Operations w/ 200M Preloaded Values

Buffered  
Binary Tree  
( $\epsilon = 0$ )

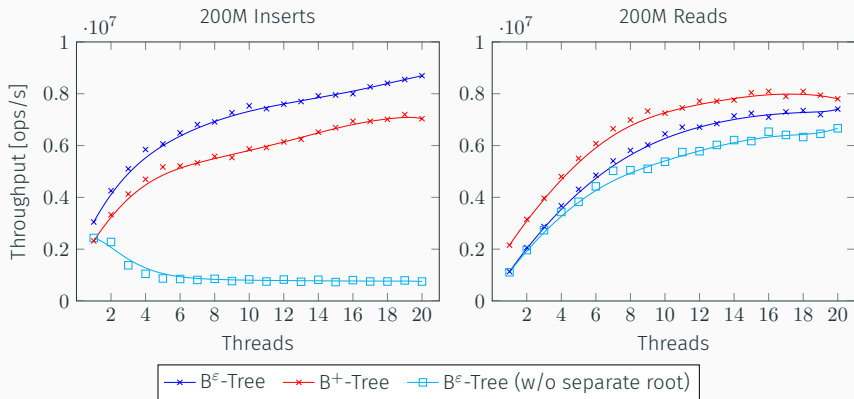


B<sup>+</sup>-Tree  
( $\epsilon = 1$ )

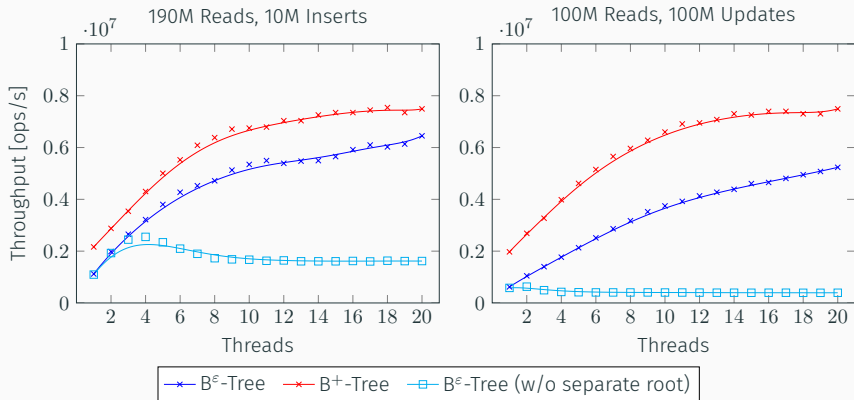


$\epsilon = 0.8 \rightarrow$  buffer slots = 109, pivots = 147

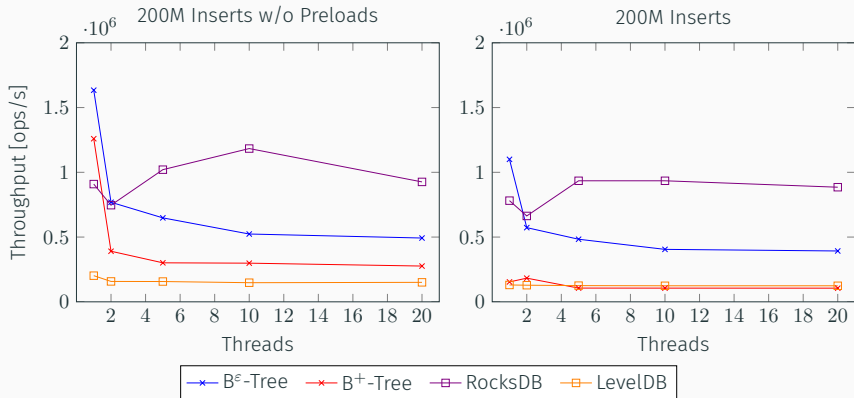
# "Optimal" Buffer | 200M Operations w/ 200M Preloaded Values



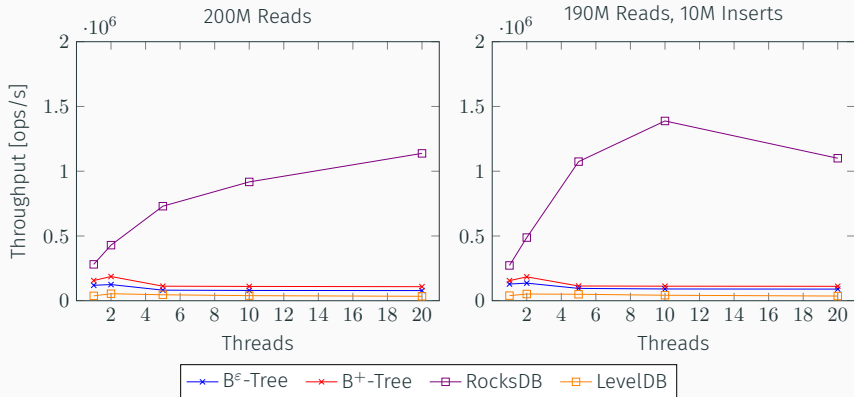
# "Optimal" Buffer | 200M Operations w/ 200M Preloaded Values



## 2Q Buffer (10GB) | 200M Operations w/ 200M Preloaded Values



## 2Q Buffer (10GB) | 200M Operations w/ 200M Preloaded Values





# $B^\varepsilon$ -Trees | Conclusion

- Asymptotically between  $B/B^+$ -Trees and LSM-Trees
- Textbook  $B^\varepsilon$ -Trees require adaption for practical use
- **Future work:**
  - Advanced page buffer
  - Optimistic Lock Coupling (OLC)
  - ...

